

MODELING CALCULATION KERNELS

Hoël Jalmin

Tutored by Arnaud Legrand and Tom Cornebize

The 21th of June, 2019

With the current need for high performance computing, and the hardware complexity:

- How to predict the duration of calculations?
- How to check if the performance is normal?

For this talk:

1. Brief presentation of the context
2. Introduction to Bayesian sampling
3. Examples of application

Modern context

- HPC systems use thousands of nodes, multiple levels of cache, hyperthreading, etc -> makes it difficult to predict performance
- Some functions are used everywhere, and called thousands of times

Polaris research

- Simulating HPL on smaller supercomputers to optimize it at a lesser cost
- Elaborated complex models but needed to evaluate and confirm the models

Examples

- The blas library is used by thousands of programs and constitutes most of HPL calculations.
- Performance variability is also caused by network communications
- Needs to check the models with bayesian sampling.

BAYESIAN STATISTICS

Model Let's say $y \sim \mathcal{N}(\mu, \sigma)$

- μ : Model **parameters**
- y : Dependent **data** (posterior)
- σ : Independent data (prior)

We observe some data and need to find model parameters


The vocabulary

- **Posterior**: The distribution of the parameters
- **Likelihood**: A function of the parameters, the model
- **Prior**: Existing knowledge of the system, guesses on the parameters values

$$\underbrace{p(\mu|y, \sigma)}_{\text{Posterior}} \propto \underbrace{p(y|\mu, \sigma)}_{\text{Likelihood}} \underbrace{p(\mu, \sigma)}_{\text{Prior}} \text{ assuming } y \sim \mathcal{M}(\mu, \sigma)$$

A BAYESIAN SAMPLER, STAN

WITH A SIMPLE EXAMPLE



Using this data, we'll try to find the parameters that were used to generate it.

THE STAN MODEL

```
library(rstan)

modelString = "data { // the observations
  int<lower=1> N; // number of points
  vector[N] x;
  vector[N] y;
}
parameters { // what we want to find
  real intercept;
  real coefficient;
  real<lower=0> sigma; // indication: sigma cannot be negative
}
model {
  // We define our priors
  intercept ~ normal(0, 10); // We know that all the parameters follow a normal
  coefficient ~ normal(0, 10);
  sigma ~ normal(0, 10);

  // Then, our likelihood function
  y ~ normal(coefficient*x + intercept, sigma);
}
"

sm = stan_model(model_code = modelString)
```

MAKING THE FIT AND CHECKING THE RESULTS

```
data = list(N=nrow(df),x=df$x,y=df$y)
fit = sampling(sm,data=data, iter=500, chains=8)
```

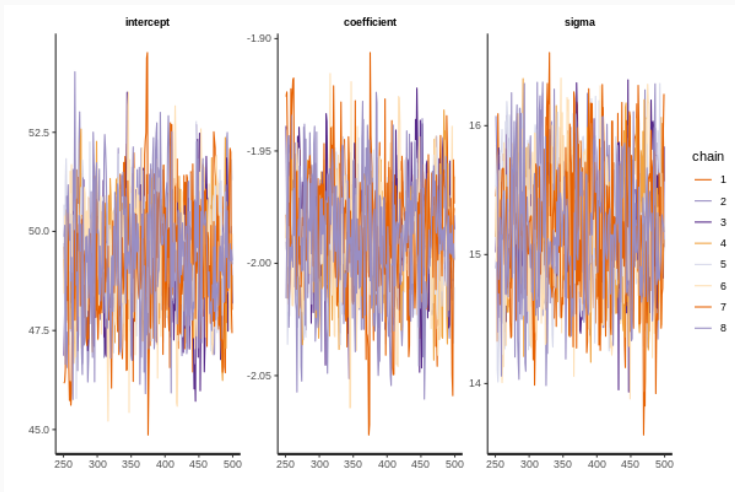
```
print(fit)
```

```
Inference for Stan model: ea4b5a288cf5f1d87215860103a9026e.
8 chains, each with iter=500; warmup=250; thin=1;
post-warmup draws per chain=250, total post-warmup draws=2000.
```

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%
intercept	49.86	0.04	1.36	47.14	48.94	49.87	50.82	52.40
coefficient	-2.00	0.00	0.02	-2.04	-2.01	-2.00	-1.98	-
1.95								
sigma	15.03	0.01	0.47	14.18	14.70	15.02	15.35	15.99
lp__	-1615.90	0.04	1.12	-1618.80	-1616.45	-1615.62	-1615.05	-1614.58
	n_eff	Rhat						
intercept	1070	1.00						
coefficient	1042	1.00						
sigma	1042	1.01						
lp__	871	1.00						

Samples were drawn using NUTS(diag_e) at Wed Jun 19 17:07:18 2019.
For each parameter, `n_eff` is a crude measure of effective sample size,
and `Rhat` is the potential scale reduction factor on split chains (at
convergence, `Rhat=1`).

CHECKING CONVERGENCE



A POLYNOMIAL MODEL WITH SPECIFIC NOISE

A MODEL DEPENDING ON THE HOST

THE IMPORTANCE OF THE PRIORS

- Possibilities**
- Modeling other calculation kernels
 - Modeling the network communications
 - Parsing and converting Stan code to C, to generate new data more efficiently
 - Anomaly detection